ME 759
High Performance Computing for Engineering Applications
Assignment 9
Due Thursday 11/10/2022 at 9:00 PM

Submit responses to all tasks which don't specify a file name to Canvas in a file called assignment9.{txt, docx, pdf, rtf, odt} (choose one of the formats). Also, all plots should be submitted in Canvas. All *source files* should be submitted in the `HW09` subdirectory on the `main` branch of your homework git repo with no subdirectories. Your `HW09` folder should have `task1.cpp`, `cluster.cpp` (new version that avoids false sharing), `task2.cpp`, `montecarlo.cpp`, and `task3.cpp`.

All commands or code must work on *Euler* without loading additional modules unless specified otherwise. Commands/code may behave differently on your computer, so be sure to test on Euler before you submit. For the OpenMP related tasks, i.e. Task 1 and Task 2, the following specifications need to be included in your `slurm` script:

- `#SBATCH --cpus-per-task=10` (or `-c 10` for short).

For Task 3 (MPI related), on top of the `slurm` script used for OpenMP, the following changes need to be made to your `slurm` script:

- Remove `#SBATCH --cpus-per-task=10` (or `-c 10` for short).

- Add `#SBATCH --ntasks-per-node=2`. This specification is only for this assignment, since only two processes are needed for task 3.

Please submit clean code. Consider using a formatter like clang-format.
* Before you begin, copy the provided files from `HW09` of the ME759 Resource Repo. Do not change any of the provided files since these files will be overwritten with clean, reference copies when grading.

---

1. (33pts) In this task, you are given a function that displays a false sharing issue. Specifically, in your code, each thread will calculate the sum of the distances between a "center point" associated with this thread and a chunk of entries of a large reference array `arr` of size `n`. The function declaration is in `cluster.h`; its definition is in `cluster.cpp`. You will need to use the techniques discussed in class (Lecture 24) to fix the false sharing issues and assess any impact of performance (false sharing vs. no false sharing). To that end:

   a) Modify the current `cluster.cpp` file to solve the false sharing issue. In this particular problem, you are allowed to modify the file `cluster.cpp`, even though it is a provided file.

   b) Write a program `task1.cpp` that will accomplish the following:

      - Create and fill with `float` type random numbers an array `arr` of length `n` where `n` is the first command line argument, see below. The range for these random floating point numbers is [0, `n`].
      - Sort `arr` (`std::sort` will do).
      - Create and fill an array `centers` of length `t`, where `t` is the number of threads and it's the second command line argument, see below. Assuming that `n` is always a multiple of `2 * t` (your code does not need to account for the cases where `n` is not a multiple of `2 * t`), the entries in the array `centers` should be (as floating point numbers):

      $$\left[\frac{n}{2t}, \frac{3n}{2t}, ..., \frac{(2t-1)n}{2t}\right].$$

      - Create and initialize with zeros an array `dists` of length `t`[1], where `t` is the number of threads and it's the second command line argument, see below.
      - Call the `cluster` function and save the output distances to the `dists` array.
      - Calculate the maximum distance in the `dists` array.
      - Print the maximum distance.
      - Print the partition ID (the thread number) that has the maximum distance.

---

[1]Note that if you choose to use padding to avoid false sharing, the size of arrays `centers` and `dists` can be changed accordingly.

- Print the time taken to run the `cluster` function in *milliseconds*.
- Compile: `g++ task1.cpp cluster.cpp -Wall -O3 -std=c++17 -o task1 -fopenmp`
- Run (where `n` is a positive integer, `t` is an integer in the range $[1, 10]$):
  `./task1 n t`
- Example expected output:
  `32516`
  `2`
  `0.032`

  A more concrete example (also given in the `cluster.h` file):
  Input: `arr` $= [0, 1, 3, 4, 6, 6, 7, 8]$, `n` $= 8$, `t` $= 2$.
  Expected `centers`: $[2, 6]$.
  Expected `dists`: $[6, 3]$, where
  $6 = |0 - 2| + |1 - 2| + |3 - 2| + |4 - 2|$,
  $3 = |6 - 6| + |6 - 6| + |7 - 6| + |8 - 6|$.
  Your code needs to output on the first line `6`, and on the second line `0`.

c) On an Euler *compute node*:

- Run `task1` for value `n` $= 5040000$, and value `t` $= 1, 2, \cdots, 10$. Generate a plot called `task1.pdf` which plots time taken by your `cluster` function vs. `t` in linear–linear scale. If there are irregular spikes in the run time, run the `cluster` function for 10 times and use the average time for plotting.

2. (33pts) In this task, you will implement an estimation of π using the Monte Carlo Method[1]. The idea is to generate `n` random floats in the range $[-r, \, r]$, where r is the radius of a circle, and count the number of floats that reside in the circle (call it `incircle`), then use `4 * incircle / n` as the estimation of π. When `n` becomes very large, this estimation could be fairly accurate. Upon implementing this method with `omp for`, you will also compare the performance difference when the `simd` directive is added.

    a) Implement in a file called `montecarlo.cpp` with the prototype specified in `montecarlo.h` the function that accomplishes the Monte Carlo Method.

    b) Your program `task2.cpp` should accomplish the following:
- Create and fill with `float`-type random numbers array `x` of length `n`, where `n` is the first command line argument, see below. `x` should be drawn from the range [`-r`, `r`], where `r` is the circle radius and it can be set to 1.0.
- Create and fill with `float`-type random numbers array `y` of length `n`, where `n` is the first command line argument, see below. `y` should be drawn from the range [`-r`, `r`], where `r` is the circle radius and it can be set to 1.0.
- Call the `montecarlo` function that returns the number of points that reside in the circle.
- Print the estimated π.
- Print the time taken to run the `montecarlo` function in *milliseconds*.
- Compile[2]: `g++ task2.cpp montecarlo.cpp -Wall -O3 -std=c++17 -o task2 -fopenmp -fno-tree-vectorize -march=native -fopt-info-vec`
- Run (where `n` is a positive integer, `t` is an integer in the range $[1, 10]$):
`./task2 n t`
- Example expected output:
`3.1416`
`0.352`

    c) On an Euler *compute node*:
- Run `task2` for $n = 10^6$, and $t = 1, 2, \cdots, 10$. Generate a figure called `task2.pdf` which includes two patterns:
    - The time taken by your `montecarlo` function **without** the `simd` directive vs. `t` in linear–linear scale.
    - The time taken by your `montecarlo` function **with** the `simd` directive vs. `t` in linear–linear scale.

    If there are irregular spikes in the run time, run the `montecarlo` function for 10 times and use the average time for plotting.

---

[1]See the details about Monte Carlo Method in this link. It also has an explanation about π estimation in the second paragraph of its Overview section

[2]Some notes about the compiler flags: `-fno-tree-vectorize` helps to avoid auto-vectorization that does not come from the OpenMP `simd` directive; `-march=native` helps to map to the vector size available on the hardware; `-fopt-info-vec` outputs a message when vectorization is achieved.)

3. (34pts) In this task, you will write an MPI program to quantify the communication latency and bandwidth between two MPI processes (executed on the same node) using `MPI_Send` and `MPI_Recv`.

a) Write a program `task3.cpp` that will accomplish the following:

- Create and fill two buffer arrays with `float`-type numbers however you like that represent the messages to be sent or received. Both arrays should have length `n`, where `n` is the first command line argument, see below.
- Initialize the variables necessary for the function call of `MPI_Send` and `MPI_Recv`.
- Implement the procedure listed below for the communication between two processes:

```
if (rank == 0) {
    // start timing t0
    MPI_Send (...);
    MPI_Recv (...);
    // end timing t0
} else if (rank == 1) {
    // start timing t1
    MPI_Recv (...);
    MPI_Send (...);
    // end timing t1
}
```

- Print the time taken for both processes to complete (`t0+t1`) in *milliseconds*. You may need to send the timing result from one rank to the other to get the total timing.
- Compile[1]: `mpicxx task3.cpp -Wall -O3 -o task3`
- Run[2] on an Euler *compute node* (meaning in a `slurm` script with proper headers), where `n` is a positive integer:
  `srun -n 2 task3 n`
- Example expected output:
  `3.45`

b) On an Euler *compute node*:

- Run `task3` for value $n = 2^1, 2^2, \cdots, 2^{25}$. Generate a plot called `task3.pdf` which plots the time taken by the communication between two processes vs. `n` in log-log scale.
- Based on the timing results, roughly estimate the latency and bandwidth of this point-to-point communication and discuss how well these estimations match with the hardware specifics.

---

[1]Use `module load mpi/mpich/4.0.2`
[2]Use `module load mpi/mpich/4.0.2`, here `srun` is equivalent to `mpirun` in the context of running an MPI application with Slurm. The `-n` flag specifies number of processes.